

Introduction to Amazon Aurora

Overview

This lab introduces you to Amazon Aurora and provides you with a basic understanding of how to use Aurora. You will follow the steps to create an Aurora instance and then connect to it.

Topics covered

After completing this lab, you will be able to:

- Create an Aurora instance
- Connect to a pre-created Amazon Elastic Compute Cloud (Amazon EC2) instance
- Configure the Amazon EC2 instance to connect to Aurora
- Query the Aurora instance

Duration

This lab requires approximately **40 minutes** to complete.

AWS service restrictions

In this lab environment, access to AWS services and service actions might be restricted to the ones that you need to complete the lab instructions. You might encounter errors if you attempt to access other services or perform actions beyond the ones that this lab describes.

Prerequisites




To successfully complete this lab, you should have some experience using the Linux operating system and have a basic understanding of structured query language (SQL).

Accessing the AWS Management Console

1. At the upper-right corner of these instructions, choose ► **Start Lab**

Troubleshooting tip: If you get an **Access Denied** error, close the error box, and choose ► **Start Lab** again.

2. The lab status can be interpreted as follows:

- A red circle next to **AWS**  at the upper-left corner of this page indicates the lab has not been started.
- A yellow circle next to **AWS**  at the upper-left corner of this page indicates the lab is starting.
- A green circle next to **AWS**  at the upper-left corner of this page indicates the lab is ready.

Wait for the lab to be ready before proceeding.

3. At the top of these instructions, choose the green circle next to **AWS** 

This option opens the AWS Management Console in a new browser tab. The system automatically sign you in.

Tip: If a new browser tab does not open, a banner or icon at the top of your browser will indicate that your browser is preventing the site

from opening pop-up windows. Choose the banner or icon, and choose **Allow pop-ups**.

4. Arrange the AWS Management Console tab so that it displays along side these instructions. Ideally, you should be able to see both browser tabs at the same time so that you can follow the lab steps.

⚠ Do not change the lab Region unless specifically instructed to do so.

It takes a few minutes to provision the resources necessary to complete this lab.

Introducing the technologies

Amazon Aurora

Aurora is a fully managed, MySQL-compatible, relational database engine that combines the performance and reliability of high-end commercial databases with the simplicity and cost-effectiveness of open-source databases. It delivers up to five times the performance of MySQL without requiring changes to most of your existing applications that use MySQL databases.

Amazon Elastic Compute Cloud (Amazon EC2)


Amazon EC2 is a web service that provides resizable compute capacity in the cloud. It is designed to make web-scale cloud computing easier for developers. Amazon EC2 reduces the time required to provision new server instances to minutes, giving you the ability to quickly scale capacity, both up and down, as your computing requirements change.

Amazon Relational Database Service (Amazon RDS)

Amazon RDS makes it easy to set up, operate, and scale a relational database in the cloud. It provides cost-efficient and resizable capacity while managing time-consuming database administration tasks, freeing you up to focus on your applications and business. Amazon RDS provides you with six database engines to choose from, including Aurora, Oracle, Microsoft SQL Server, PostgreSQL, MySQL, and MariaDB.

Task 1: Create an Aurora instance

In this task, you create an Aurora database (DB) instance.

5. In the AWS Management Console, choose the  **Services** menu. Choose **Database**, and choose **RDS**.
6. In the left navigation menu, choose **Databases**.
7. Choose **Create database** and then configure the following options:
 - For **Choose a database creation method**, choose **Standard create**.
 - For **Engine options**, choose **Amazon Aurora**.
 - In the **Engine options** section for **Replication features**, choose **Single-master** (if this option is not selected by default).
 - For **Templates**, choose **Dev/Test**.
8. In the **Settings** section, configure the following options:
 - **DB cluster identifier**: Enter `aurora`
 - **Master username**: Enter `admin`
 - **Master password**: Enter `admin123`
 - **Confirm password**: Enter `admin123`
9. In the **DB instance class** section, choose **Burstable classes**, and choose **db.t3.small** from the dropdown list.
10. In the **Availability & durability** section for **Multi-AZ deployment**, choose **Don't create an Aurora Replica**.

i Amazon RDS Multi-AZ deployments provide enhanced availability and durability for DB instances, making them a natural fit for production database workloads. When you provision a Multi-AZ DB instance, Amazon RDS automatically creates a primary DB instance and synchronously replicates the data to a standby instance in a different Availability Zone.

Because this is a lab environment, you do not need to perform a multi-AZ deployment.

11. In the **Connectivity** section, configure the following options:

- For **Virtual private cloud (VPC)**, choose **LabVPC**.
- For **Subnet group**, choose **dbsubnetgroup**.
- For **Public access**, choose **No**.
- For **VPC security group**, choose **Choose existing**.
- For **Existing VPC security groups**, remove the **default** security group.
- From the **Existing VPC security groups** dropdown list, choose **DBSecurityGroup**.

i Subnets are segments of a virtual private cloud (VPC) IP address range that you designate to group your resources based on security and operational needs. A DB subnet group is a collection of subnets (typically private) that you create in a VPC and that you then designate for your DB instances. With a DB subnet group, you can specify a particular VPC when creating DB instances using the command line interface (CLI) or application programming interface (API); if you use the console, you can select the VPC and subnets that you want to use.

The **aurora** subnet group was created for you when you launched the lab using AWS CloudFormation.

i You can use the Amazon Virtual Private Cloud (Amazon VPC) service to launch AWS resources into a virtual network that you've defined. This virtual network closely resembles a traditional network that you'd operate in your own data center, with the benefits of using the scalable infrastructure of AWS.

12. Expand **Additional configuration**, and for **Initial database name**, enter `world`

13. In the **Encryption** section, clear the check box for **Enable encryption**.

i You can encrypt your Amazon RDS instances and snapshots at rest by enabling the encryption option for your RDS DB instance. Data that is encrypted at rest includes the underlying storage for a DB instance, its automated backups, read replicas, and snapshots.

14. In the **Monitoring** section, clear the check box for **Enable Enhanced monitoring**.

15. In the **Maintenance** section, clear the check box for **Enable auto minor version upgrade**.

16. Scroll to the bottom of the screen, and then choose **Create database**

Your Aurora DB instance is in the process of launching and can take up to 5 minutes to launch. However, you can continue to the next task.

Task 2: Connect to an Amazon EC2 Linux instance

In this task, you log into to your Amazon EC2 Linux instance. This instance was launched for you when you started your lab using CloudFormation.

17. In the AWS Management Console, choose the **Services** menu. Choose **Compute**, and then choose **EC2**.

18. In the left navigation menu, choose **Instances**.

19. Next to the instance labelled **Command Host**, select the check box, and then choose **Connect**.

Note: If you do not see the **Command Host** instance, the lab is possibly still being provisioned, or you may be using another Region.

20. For **Connect to instance**, choose **Session Manager**.

21. Choose **Connect** to open a terminal window.

Note: If the **Connect** button is not available, wait for a few minutes and try again.



Task 3: Configure the Amazon EC2 Linux instance to connect to Aurora

In this task, you configure the Amazon EC2 Linux instance to connect to Aurora.

22. To configure the instance with the **MariaDB client**, run the following command. The MariaDB client is used to connect to the Aurora

instance that you just created.

```
sudo yum install mariadb -y
```

23. Using a different browser tab, go back to the AWS Management Console. Choose the  **Services** menu. Choose **Database**, and then choose **RDS**.
24. In the left navigation menu, choose **Databases**.
25. Wait for **aurora-instance-1** to display  **Available**.
26. Choose **aurora**.
27. Choose the **Connectivity & security** tab, and in the **Endpoints** section, copy the **Endpoint name** for the **Writer instance** to your text editor.

The endpoint should look similar to the following: **aurora.cluster-cabcdefghijklm.us-west-2.rds.amazonaws.com**

Types of Aurora endpoints

An endpoint is represented as an Aurora specific URL that contains a host address and a port. The following types of endpoints are available from an Aurora DB cluster.

- **Cluster endpoint:**

A cluster endpoint for an Aurora DB cluster connects to the current primary DB instance for that DB cluster. This endpoint is the only one that can perform write operations such as DDL statements. Because of this, the cluster endpoint is the one that you connect to when you first set up a cluster or when your cluster contains only a single DB instance.

Each Aurora DB cluster has one cluster endpoint and one primary DB instance.

You use the cluster endpoint for all write operations on the DB cluster, including inserts, updates, deletes, and DDL changes. You can also use the cluster endpoint for read operations, such as queries.

The cluster endpoint provides failover support for read/write connections to the DB cluster. If the current primary DB instance of a DB cluster fails, Aurora automatically fails over to a new primary DB instance. During a failover, the DB cluster continues to serve connection requests to the cluster endpoint from the new primary DB instance, with minimal interruption of service.

The following example illustrates a cluster endpoint for an Aurora MySQL DB cluster.

mydbcluster.cluster-123456789012.us-west-2.rds.amazonaws.com:3306

- **Reader endpoint:**

A reader endpoint for an Aurora DB cluster connects to one of the available Aurora replicas for that DB cluster. Each Aurora DB cluster has one reader endpoint. If there is more than one Aurora replica, the reader endpoint directs each connection request to one of the Aurora replicas.

The reader endpoint provides load-balancing support for read-only connections to the DB cluster. Use the reader endpoint for read operations, such as queries. You can't use the reader endpoint for write operations.

The DB cluster distributes connection requests to the reader endpoint among the available Aurora replicas. If the DB cluster contains only a primary DB instance, the reader endpoint serves connection requests from the primary DB instance. If one or more Aurora replicas are created for that DB cluster, subsequent connections to the reader endpoint are load balanced among the replicas.

The following example illustrates a reader endpoint for an Aurora MySQL DB cluster.

mydbcluster.cluster-ro-123456789012.us-west-2.rds.amazonaws.com:3306

28. In the following command, replace *<endpoint_goes_here>* with the endpoint that you copied to your text editor.

```
mysql -u admin --password='admin123' -h <endpoint_goes_here>
```

Your command should look similar to the following:

```
mysql -u admin --password='admin123' -h mydbcluster.cluster-123456789012.us-west-2.rds.amazonaws.com
```

The MySQL Command-Line Client is a SQL shell which enables interaction with database engines. More information is available [here](#).

Switch	Description
-u or --user	The MySQL username used to connect to a database instance.
-p or --password	The MySQL password used to connect to a database instance.
-h or --host	The host address of the database engine.

Copy the command to your clipboard.

29. Return to the **Session Manager** browser tab that was used to connect to the **Command Host**. To connect to the Aurora instance, run the command you had copied in the previous step.

Task 4: Create a table and insert and query records

In this task, you learn how to create a table in a database, load data, and run a query.

30. To list the available databases, run the following command.

```
SHOW DATABASES;
```

31. To switch to the **world** database that you created in Task 1 when you provisioned the Aurora instance, run the following command.

```
USE world;
```

32. To create a new table in the **world** database, run the following command.

```
CREATE TABLE `country` (  
  `Code` CHAR(3) NOT NULL DEFAULT '',  
  `Name` CHAR(52) NOT NULL DEFAULT '',  
  `Continent` enum('Asia','Europe','North America','Africa','Oceania','Antarctica','South America') NOT NULL  
  DEFAULT 'Asia',  
  `Region` CHAR(26) NOT NULL DEFAULT '',  
  `SurfaceArea` FLOAT(10,2) NOT NULL DEFAULT '0.00',  
  `IndepYear` SMALLINT(6) DEFAULT NULL,  
  `Population` INT(11) NOT NULL DEFAULT '0',  
  `LifeExpectancy` FLOAT(3,1) DEFAULT NULL,  
  `GNP` FLOAT(10,2) DEFAULT NULL,  
  `GNP01d` FLOAT(10,2) DEFAULT NULL,  
  `LocalName` CHAR(45) NOT NULL DEFAULT '',  
  `GovernmentForm` CHAR(45) NOT NULL DEFAULT '',  
  `Capital` INT(11) DEFAULT NULL,  
  `Code2` CHAR(2) NOT NULL DEFAULT '',  
  PRIMARY KEY (`Code`)  
);
```

33. To insert new records into the **country** table that you just created, run the following commands.

```

INSERT INTO `country` VALUES ('GAB','Gabon','Africa','Central
Africa',267668.00,1960,1226000,50.1,5493.00,5279.00,'Le Gabon','Republic',902,'GA');

INSERT INTO `country` VALUES ('IRL','Ireland','Europe','British
Islands',70273.00,1921,3775100,76.8,75921.00,73132.00,'Ireland/Éire','Republic',1447,'IE');

INSERT INTO `country` VALUES ('THA','Thailand','Asia','Southeast
Asia',513115.00,1350,61399000,68.6,116416.00,153907.00,'Prathet Thai','Constitutional Monarchy',3320,'TH');

INSERT INTO `country` VALUES ('CRI','Costa Rica','North America','Central
America',51100.00,1821,4023000,75.8,10226.00,9757.00,'Costa Rica','Republic',584,'CR');

INSERT INTO `country` VALUES ('AUS','Australia','Oceania','Australia and New
Zealand',7741220.00,1901,18886000,79.8,351182.00,392911.00,'Australia','Constitutional Monarchy,
Federation',135,'AU');

```

34. To query the table, run the following **SELECT** statement.

```
SELECT * FROM country WHERE GNP > 35000 and Population > 10000000;
```

The query should return two records.

Conclusion

👏 Congratulations! You have now successfully:

- Created an Aurora instance
- Connected to a pre-created Amazon Elastic Compute Cloud (Amazon EC2) instance
- Configured the Amazon EC2 instance to connect to Aurora
- Queried the Aurora instance

Lab complete 🎓

35. Choose **End Lab** at the top of this page, and then select **Yes** to confirm that you want to end the lab.

36. An **Ended AWS Lab Successfully** message is briefly displayed indicating that the lab has ended.